

Ashwani Kansal

+91-9996659688 | ashwanikansal.dev@gmail.com | LinkedIn Profile | Github Profile

PROFESSIONAL SUMMARY

AI Engineer specializing in agentic AI, LLM application development, and Generative AI systems, building production-grade RAG pipelines, multi-agent workflows (LangChain, LangGraph, MCP), and LLM-powered automation using Python, Anthropic Claude, and OpenAI APIs. Experienced across the full AI delivery lifecycle, from data pipelines and retrieval systems to agentic tool integration and scalable API-based deployment for enterprise and product use cases.

EDUCATION

National Institute of Technology, Kurukshetra
Masters in Computer Applications / Overall CGPA - 9.0

Kurukshetra, India
2021 – 2024

EXPERIENCE

AI/LLM Application Engineer

Pride Global

May 2026 – Present

Bengaluru (Hybrid)

- Building an onboarding automation platform using LangGraph-based agentic workflows to orchestrate document processing and system integrations across JobDiva, Microsoft Graph API, and MongoDB, reducing manual onboarding effort.
- Integrating Microsoft Azure Document Intelligence for automated extraction and validation of onboarding documents, improving data accuracy and processing speed.
- Accelerating agentic workflow development using Claude Code and the Anthropic API, enabling rapid iteration on LLM-driven automation.
- Mentoring intern engineers building RAG-based document Q&A tools on local LLMs served via Ollama, enabling privacy-preserving, on-prem inference.
- Guided an intern-led MCP (Model Context Protocol) server project that let users upload documents and query them through the Claude Desktop app as the MCP client.

AI Engineer

Genpact

Oct 2024 – Dec 2025

Bengaluru

- Built an LLM-powered document intelligence pipeline using OpenAI, LangChain, and embeddings, enabling semantic deduplication and consolidation across 20–50 unstructured documents and reducing manual review time.
- Designed RAG pipelines with FAISS and transformer-based models, improving semantic retrieval accuracy and enabling reliable Q&A over unstructured enterprise data.
- Developed FastAPI microservices for document ingestion, chunking, vector search, and LLM inference, enabling scalable reuse of GenAI capabilities across internal teams.
- Optimized PySpark transformations and Databricks jobs, improving pipeline reliability and governance using Delta Tables and Unity Catalog.

AI Engineer Intern

Genpact

Feb 2024 – Jul 2024

Bengaluru

- Designed an AWS-based healthcare data lake (S3) for structured and unstructured medical data, enabling scalable ingestion for downstream ML applications.
- Built ETL workflows using Apache Airflow to automate ingestion and preprocessing, reducing manual pipeline maintenance.
- Integrated MLflow for experiment tracking and model evaluation, improving reproducibility across ML workflows.

SKILLS

Agentic AI Frameworks: LangChain, LangGraph, Multi-Agent Systems, MCP (Model Context Protocol), Claude Code

GenAI & NLP: RAG Pipelines, OpenAI API, Anthropic API, Embeddings, FAISS, Ollama, Prompt Engineering

Machine Learning & Data: scikit-learn, TensorFlow/Keras, MLflow, Feature Engineering, Model Evaluation

Data Engineering: FastAPI, Python, PySpark, SparkSQL, Databricks, Delta Tables, Airflow

Integrations & APIs: MongoDB, Microsoft Graph API, Azure Document Intelligence

Cloud & DevOps: AWS (S3, EC2, Lambda), Docker, Git, SQL, NoSQL, Vector Stores, CI/CD